

# Virtual Memory

## Chapter 8

# Hardware and Control Structures

- Memory references are dynamically translated into physical addresses at run time
  - A process may be swapped in and out of main memory such that it occupies different regions
- A process may be broken up into pieces that do not need to be located contiguously in main memory
  - All pieces of a process do not need to be loaded in main memory during execution



# Execution of a Program

- Operating system brings into main memory a few pieces of the program
- Resident set - portion of process that is in main memory
- An interrupt is generated when an address is needed that is not in main memory
- Operating system places the process in a blocking state



# Execution of a Program

- Piece of process that contains the logical address is brought into main memory
  - Operating system issues a disk I/O Read request
  - Another process is dispatched to run while the disk I/O takes place
  - An interrupt is issued when disk I/O complete which causes the operating system to place the affected process in the Ready state



# Advantages of Breaking up a Process

- More processes may be maintained in main memory
  - Only load in some of the pieces of each process
  - With so many processes in main memory, it is very likely a process will be in the Ready state at any particular time
- A process may be larger than all of main memory



# Types of Memory

- Real memory
  - Main memory
- Virtual memory
  - Memory on disk
  - Allows for effective multiprogramming and relieves the user of tight constraints of main memory



# Thrashing

- Swapping out a piece of a process just before that piece is needed
- The processor spends most of its time swapping pieces rather than executing user instructions



# Principle of Locality

- Program and data references within a process tend to cluster
- Only a few pieces of a process will be needed over a short period of time
- Possible to make intelligent guesses about which pieces will be needed in the future
- This suggests that virtual memory may work efficiently





# Support Needed for Virtual Memory

- Hardware must support paging and segmentation
- Operating system must be able to management the movement of pages and/or segments between secondary memory and main memory



# Paging

- Each process has its own page table
- Each page table entry contains the frame number of the corresponding page in main memory
- A bit is needed to indicate whether the page is in main memory or not




# Modify Bit in Page Table

- Another modify bit is needed to indicate if the page has been altered since it was last loaded into main memory
- If no change has been made, the page does not have to be written to the disk when it needs to be swapped out



# Page Table Entries

Virtual Address



Page Table Entry



(a) Paging only

Figure 8.2 Typical Memory Management Formats



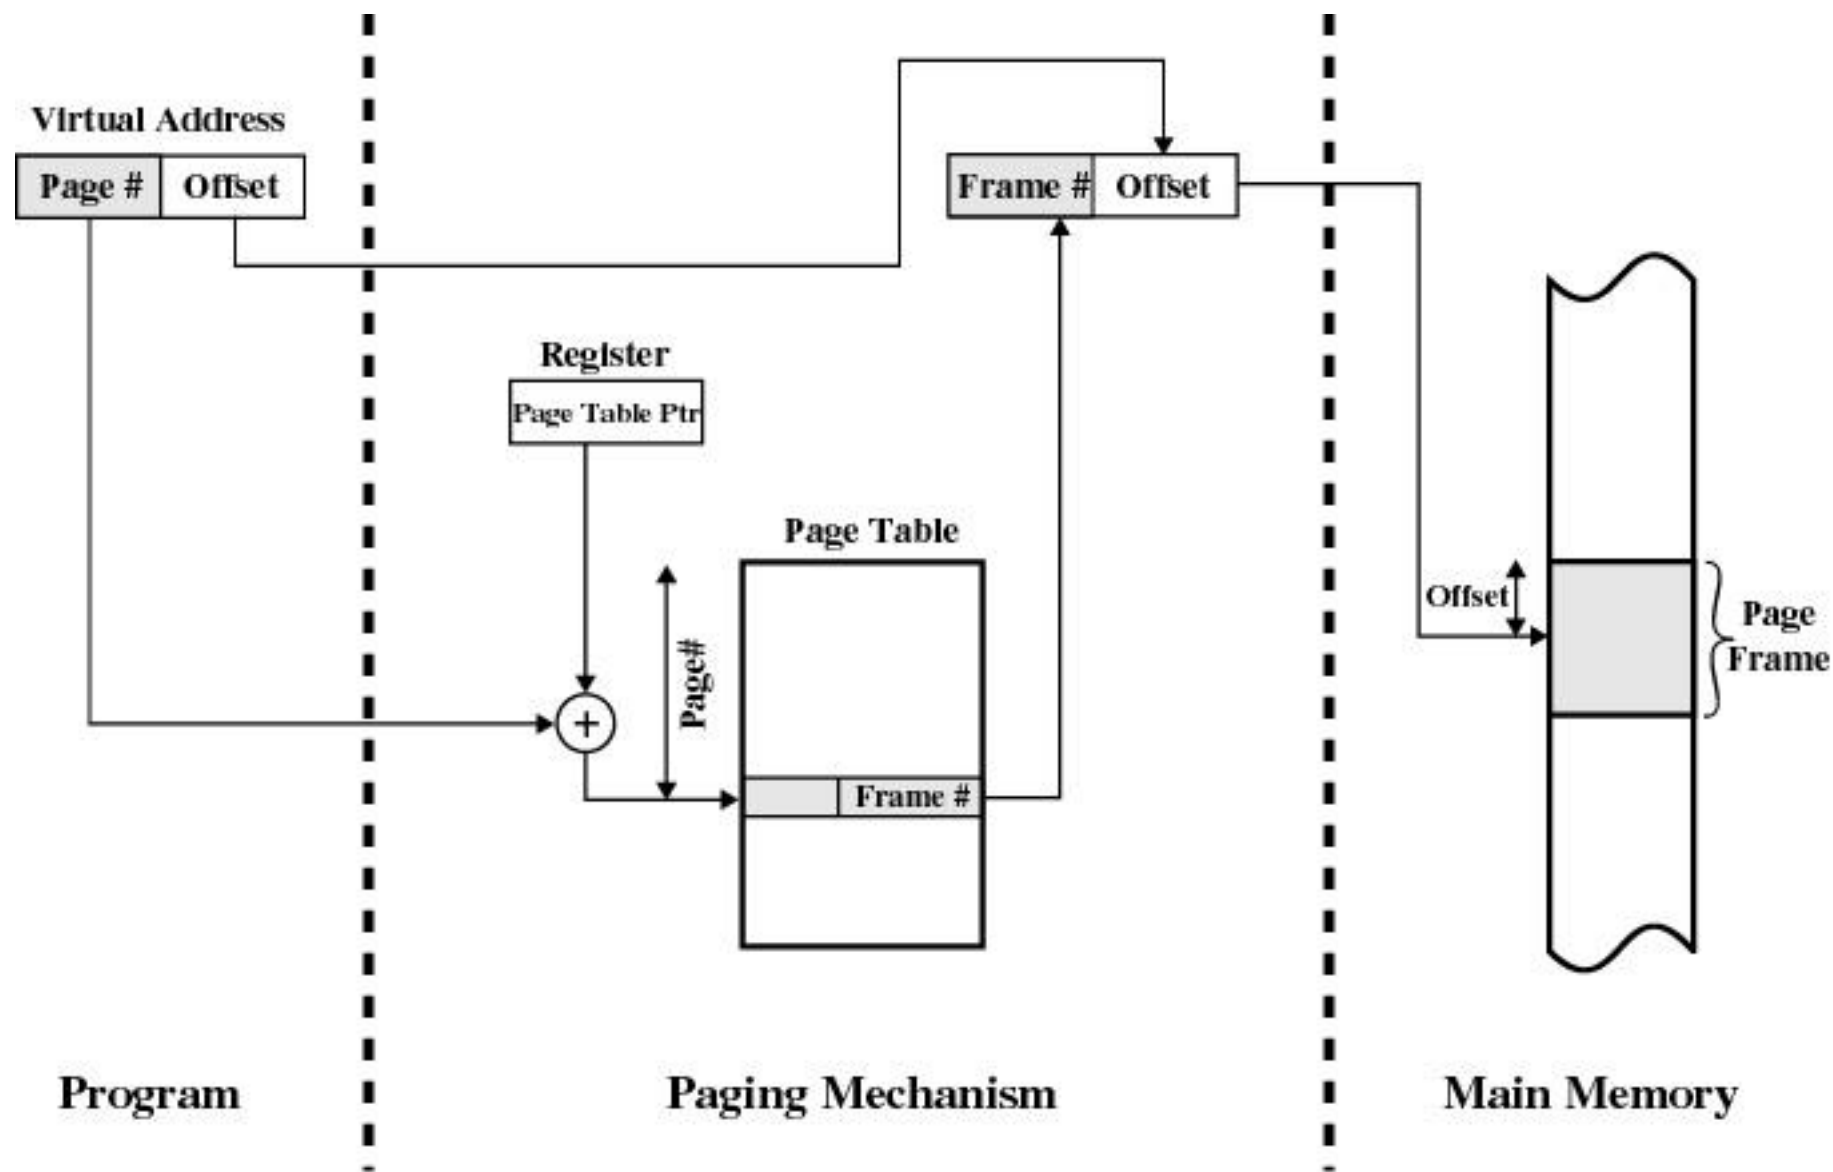


Figure 8.3 Address Translation in a Paging System

# Two-Level Scheme for 32-bit Address




Figure 8.4 A Two-Level Hierarchical Page Table [JACO98a]



# Page Tables

- The entire page table may take up too much main memory
- Page tables are also stored in virtual memory
- When a process is running, part of its page table is in main memory



# Translation Lookaside Buffer

- Each virtual memory reference can cause two physical memory accesses
  - one to fetch the page table entry
  - one to fetch the data
- To overcome this problem a high-speed cache is set up for page table entries
  - called the TLB - Translation Lookaside Buffer





# Translation Lookaside Buffer

- Contains page table entries that have been most recently used
- Functions same way as a memory cache



# Translation Lookaside Buffer

- Given a virtual address, processor examines the TLB
- If page table entry is present (a hit), the frame number is retrieved and the real address is formed
- If page table entry is not found in the TLB (a miss), the page number is used to index the process page table



# Translation Lookaside Buffer

- First checks if page is already in main memory
  - if not in main memory a page fault is issued
- The TLB is updated to include the new page entry



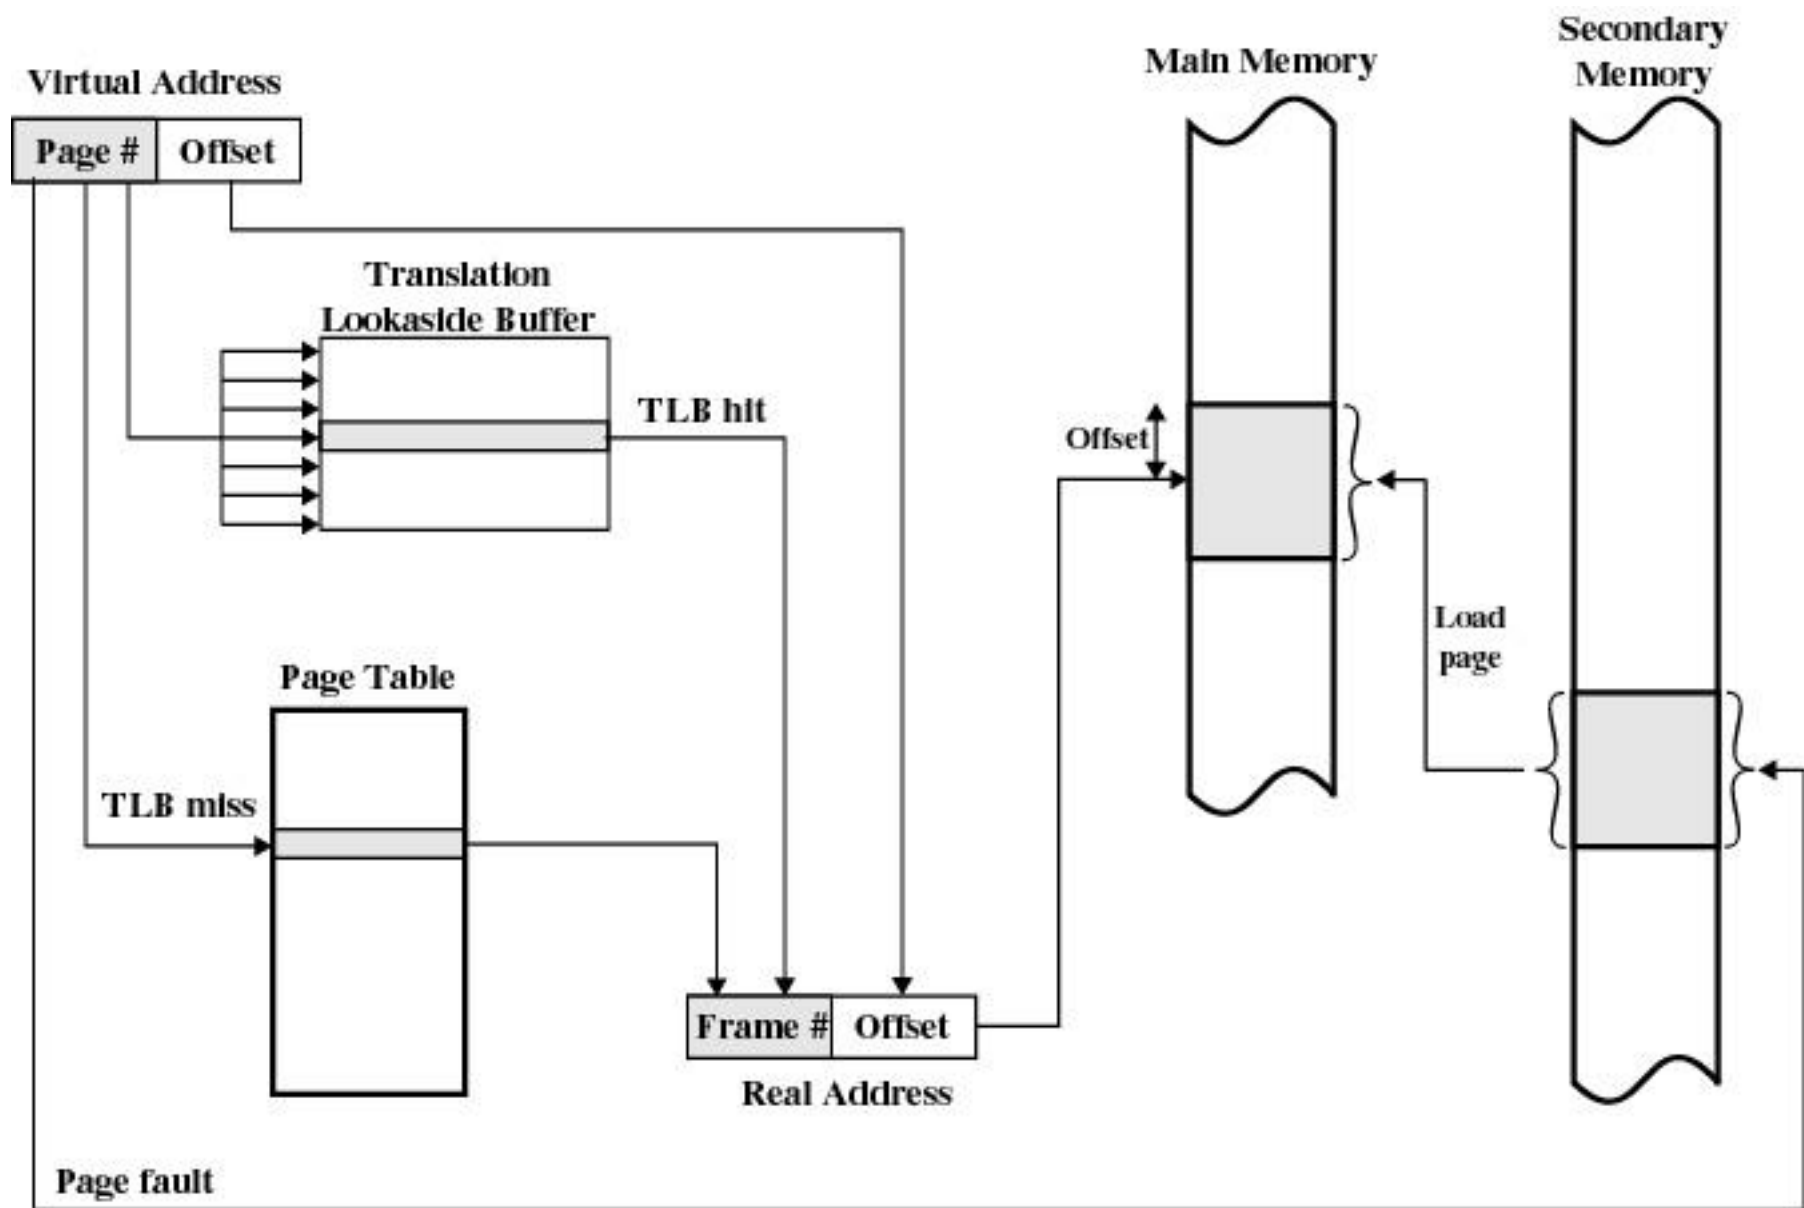


Figure 8.7 Use of a Translation Lookaside Buffer

# Page Size


- Smaller page size, less amount of internal fragmentation
- Smaller page size, more pages required per process
- More pages per process means larger page tables
- Larger page tables means large portion of page tables in virtual memory
- Secondary memory is designed to efficiently transfer large blocks of data so a large page size is better




# Page Size

- Small page size, large number of pages will be found in main memory
- As time goes on during execution, the pages in memory will all contain portions of the process near recent references. Page faults low.
- Increased page size causes pages to contain locations further from any recent reference. Page faults rise.





(a) Page Size



(b) Number of Page Frames Allocated

$P$  = size of entire process

$W$  = working set size

$N$  = total number of pages in process

**Figure 8.11 Typical Paging Behavior of a Program**

# Example Page Sizes

**Table 8.2 Example Page Sizes**

| Computer               | Page Size             |
|------------------------|-----------------------|
| Atlas                  | 512 48-bit words      |
| Honeywell-Multics      | 1024 36-bit word      |
| IBM 370/XA and 370/ESA | 4 Kbytes              |
| VAX family             | 512 bytes             |
| IBM AS/400             | 512 bytes             |
| DEC Alpha              | 8 Kbytes              |
| MIPS                   | 4 kbytes to 16 Mbytes |
| UltraSPARC             | 8 Kbytes to 4 Mbytes  |
| Pentium                | 4 Kbytes or 4 Mbytes  |
| PowerPc                | 4 Kbytes              |





# Segmentation

- May be unequal, dynamic size
- Simplifies handling of growing data structures
- Allows programs to be altered and recompiled independently
- Lends itself to sharing data among processes
- Lends itself to protection



# Segment Tables

- corresponding segment in main memory
- Each entry contains the length of the segment
- A bit is needed to determine if segment is already in main memory
- Another bit is needed to determine if the segment has been modified since it was loaded in main memory



# Segment Table Entries

Virtual Address



Segment Table Entry



(b) Segmentation only

**Figure 8.2 Typical Memory Management Formats**




# Combined Paging and Segmentation

- Paging is transparent to the programmer
- Paging eliminates external fragmentation
- Segmentation is visible to the programmer
- Segmentation allows for growing data structures, modularity, and support for sharing and protection
- Each segment is broken into fixed-size pages



# Combined Segmentation and Paging

Virtual Address



Segment Table Entry



Page Table Entry



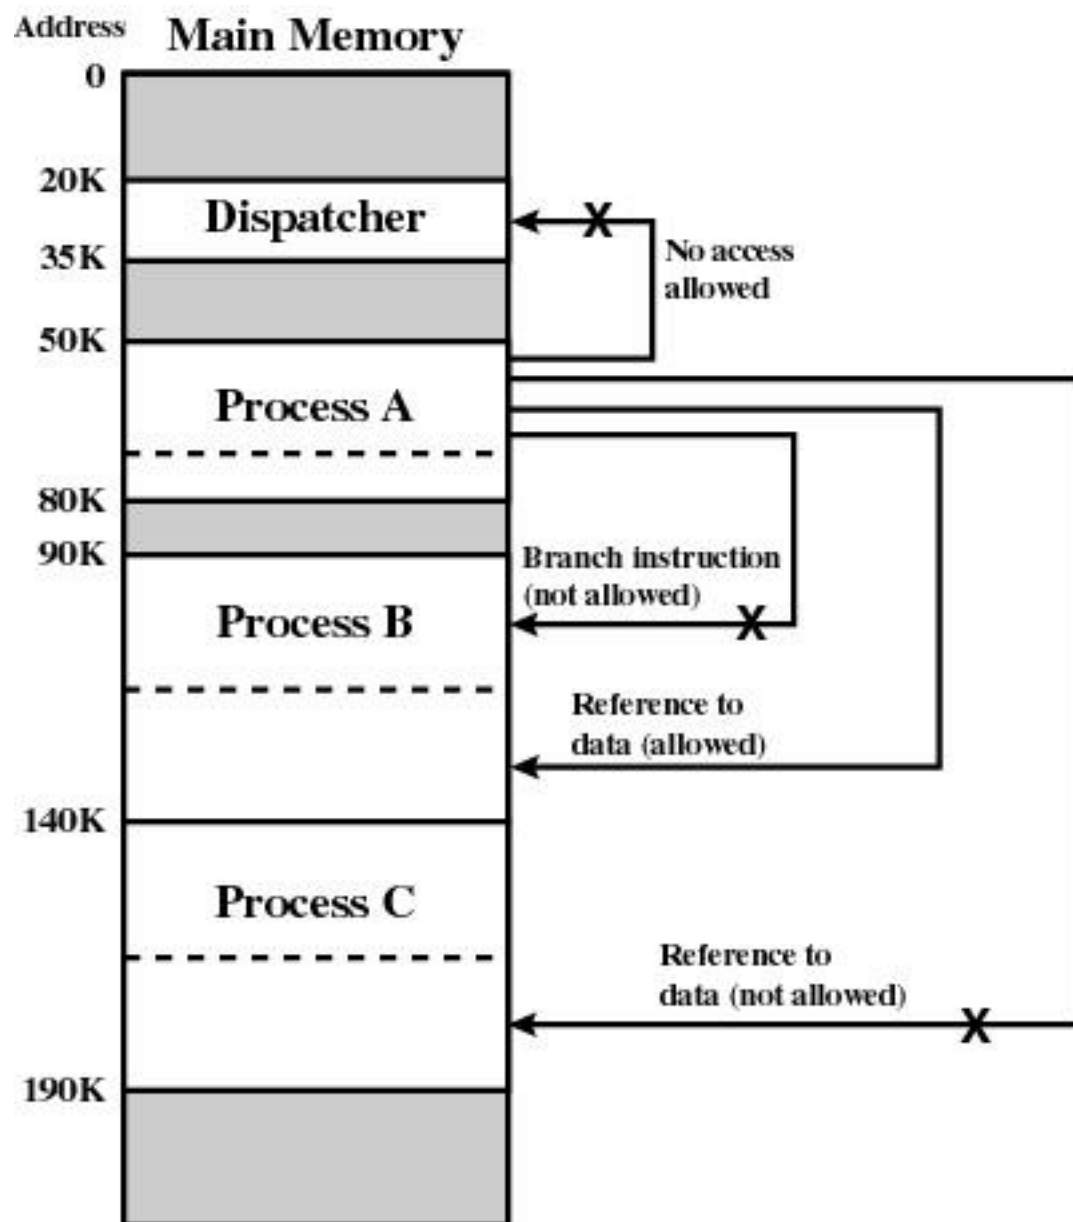
P= present bit

M = Modified bit

(c) Combined segmentation and paging

Figure 8.2 Typical Memory Management Formats





**Figure 8.14 Protection Relationships Between Segments**

# Fetch Policy

- Fetch Policy
  - Determines when a page should be brought into memory
  - Demand paging only brings pages into main memory when a reference is made to a location on the page
    - Many page faults when process first started
  - Prepaging brings in more pages than needed
    - More efficient to bring in pages that reside contiguously on the disk



# Replacement Policy

- Placement Policy
  - Which page is replaced?
  - Page removed should be the page least likely to be referenced in the near future
  - Most policies predict the future behavior on the basis of past behavior





# Replacement Policy

- Frame Locking
  - If frame is locked, it may not be replaced
  - Kernel of the operating system
  - Control structures
  - I/O buffers
  - Associate a lock bit with each frame



# Basic Replacement Algorithms

- Optimal policy
  - Selects for replacement that page for which the time to the next reference is the longest
  - Impossible to have perfect knowledge of future events



# Basic Replacement Algorithms

- Least Recently Used (LRU)
  - Replaces the page that has not been referenced for the longest time
  - By the principle of locality, this should be the page least likely to be referenced in the near future
  - Each page could be tagged with the time of last reference. This would require a great deal of overhead.



# Basic Replacement Algorithms


- First-in, first-out (FIFO)
  - Treats page frames allocated to a process as a circular buffer
  - Pages are removed in round-robin style
  - Simplest replacement policy to implement
  - Page that has been in memory the longest is replaced
  - These pages may be needed again very soon



# Basic Replacement Algorithms


- Clock Policy
  - Additional bit called a use bit
  - When a page is first loaded in memory, the use bit is set to 0
  - When the page is referenced, the use bit is set to 1
  - When it is time to replace a page, the first frame encountered with the use bit set to 0 is replaced.
  - During the search for replacement, each use bit set to 1 is changed to 0





(a) State of buffer just prior to a page replacement

**Figure 8.16 Example of Clock Policy Operation**



(b) State of buffer just after the next page replacement

**Figure 8.16 Example of Clock Policy Operation**

# Cleaning Policy

- Demand cleaning
  - a page is written out only when it has been selected for replacement
- Precleaning
  - pages are written out in batches





# Cleaning Policy

- Best approach uses page buffering
  - Replaced pages are placed in two lists
    - Modified and unmodified
  - Pages in the modified list are periodically written out in batches
  - Pages in the unmodified list are either reclaimed if referenced again or lost when its frame is assigned to another page



# Load Control

- Determines the number of processes that will be resident in main memory
- Too few processes, many occasions when all processes will be blocked and much time will be spent in swapping
- Too many processes will lead to thrashing



# Reducing Load by Process Suspension

- Lowest priority process is suspended
- Faulting process is suspended
  - this process does not have its working set in main memory so it will be blocked anyway
- Last process activated
  - this process is least likely to have its working set resident



# Process Suspension

- Process with smallest resident set
  - this process requires the least future effort to reload
- Largest process
  - obtains the most free frames
- Process with the largest remaining execution window



# UNIX and Solaris Memory Management

- Paging System
  - Page table: one per process
  - Disk block descriptor: disk copy of a page
  - Page frame data table: frame-page mapping
  - Swap-use table: one per swap device



# Data Structures

|                   |     |               |        |           |       |         |
|-------------------|-----|---------------|--------|-----------|-------|---------|
| Page frame number | Age | Copy on write | Modify | Reference | Valid | Protect |
|-------------------|-----|---------------|--------|-----------|-------|---------|

(a) Page table entry

|                    |                     |                 |
|--------------------|---------------------|-----------------|
| Swap device number | Device block number | Type of storage |
|--------------------|---------------------|-----------------|

(b) Disk block descriptor

Figure 8.22 UNIX SVR4 Memory Management Formats



# Data Structures

|            |                 |                |              |                |
|------------|-----------------|----------------|--------------|----------------|
| Page state | Reference count | Logical device | Block number | Pfdata pointer |
|------------|-----------------|----------------|--------------|----------------|

(c) Page frame data table entry

|                 |                          |
|-----------------|--------------------------|
| Reference count | Page/storage unit number |
|-----------------|--------------------------|

(d) Swap-use table entry

Figure 8.22 UNIX SVR4 Memory Management Formats





# UNIX and Solaris Memory Management

- Page Replacement
  - refinement of the clock policy
  - First sweep sets use bit to 0
  - After some time, second sweep checks use bits, if still zero, this page can be replaced
- Kernel Memory Allocator
  - most blocks are smaller than a typical page size so buddy system is used and paging is not used

